

Deep Learning Theory and Practice

Lecture 13

Convolutional Neural Network Architectures

Dr. Ted Willke
willke@pdx.edu

Monday, May 13, 2019

Review of Lecture 12

- **Parameter tying and sharing**

- Take advantage of when you know there are dependencies between parameters

Parameter tying and parameter sharing

Toy example: Two models performing same classification task (but with slightly different input and output distributions)

Model A with $W^{(A)}$

$$\hat{y}^{(A)} = h(W^{(A)}, \mathbf{x})$$

Model B with $W^{(B)}$

$$\hat{y}^{(B)} = h(W^{(B)}, \mathbf{x})$$

Similar enough that $\forall i, w_i^{(A)}$ should be close to $w_i^{(B)}$

Can leverage this through regularization:

$$\Omega(W^{(A)}, W^{(B)}) = \|W^{(A)} - W^{(B)}\|_F^2 \quad (\text{parameter tying through norm penalty})$$

8


- Additional benefits if parameters can be shared.

CNNs save memory and computation this way.

- **Convolutional neural networks**

- Utilizes spatial organization (receptive fields)


Convolutional layer



26

Review of Lecture 12

- Utilizes hierarchical organization



- Use any number of filters/activation maps

What filters learn from images

32 5x5 filters

Activations:

one filter = one activation map (if convolved over image)

Activation:

Convolutional layers are **convolutional** because we are convolving a filter (kernel) over an input (the 'image').

Most libraries don't flip the kernel, so it's really the **cross-correlation** that we're computing.

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

element-wise multiplication and sum of the kernel and the image

31

Review of Lecture 12

- Zero padding the border to retain dims

Zero padding the border

Common in practice to zero pad the border (boundary condition).

0	0	0	0	0	0	0	0	0
0	0	0						0
0	0	0						0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

For example, input is 7x7, filter is 3x3, applied with **stride 1**.

If we **pad the border** with 1 pixel, what is the output dim?

7x7 output!

Common to see **Conv** layers with stride of 1, filters of size $F \times F$, and zero-padding of $(F - 1) / 2$.

Will preserve size spatially:

$F = 3 \longrightarrow$ zero pad with 1.

$F = 5 \longrightarrow$ zero pad with 2.

$F = 7 \longrightarrow$ zero pad with 3.


46

- CNNs reduce number of parameters!

Examples

If the input volume is 32x32x3, what is the **number of parameters** in this layer?

10 5x5 filters with stride of 1, pad 2?



Number of parameters:

Each filter has $5 \times 5 \times 3 + 1 = 76$ parameters
(+1 for the bias)

$\longrightarrow 76 \times 10 = 760!$

(much fewer than a FC layer with this input size)

49

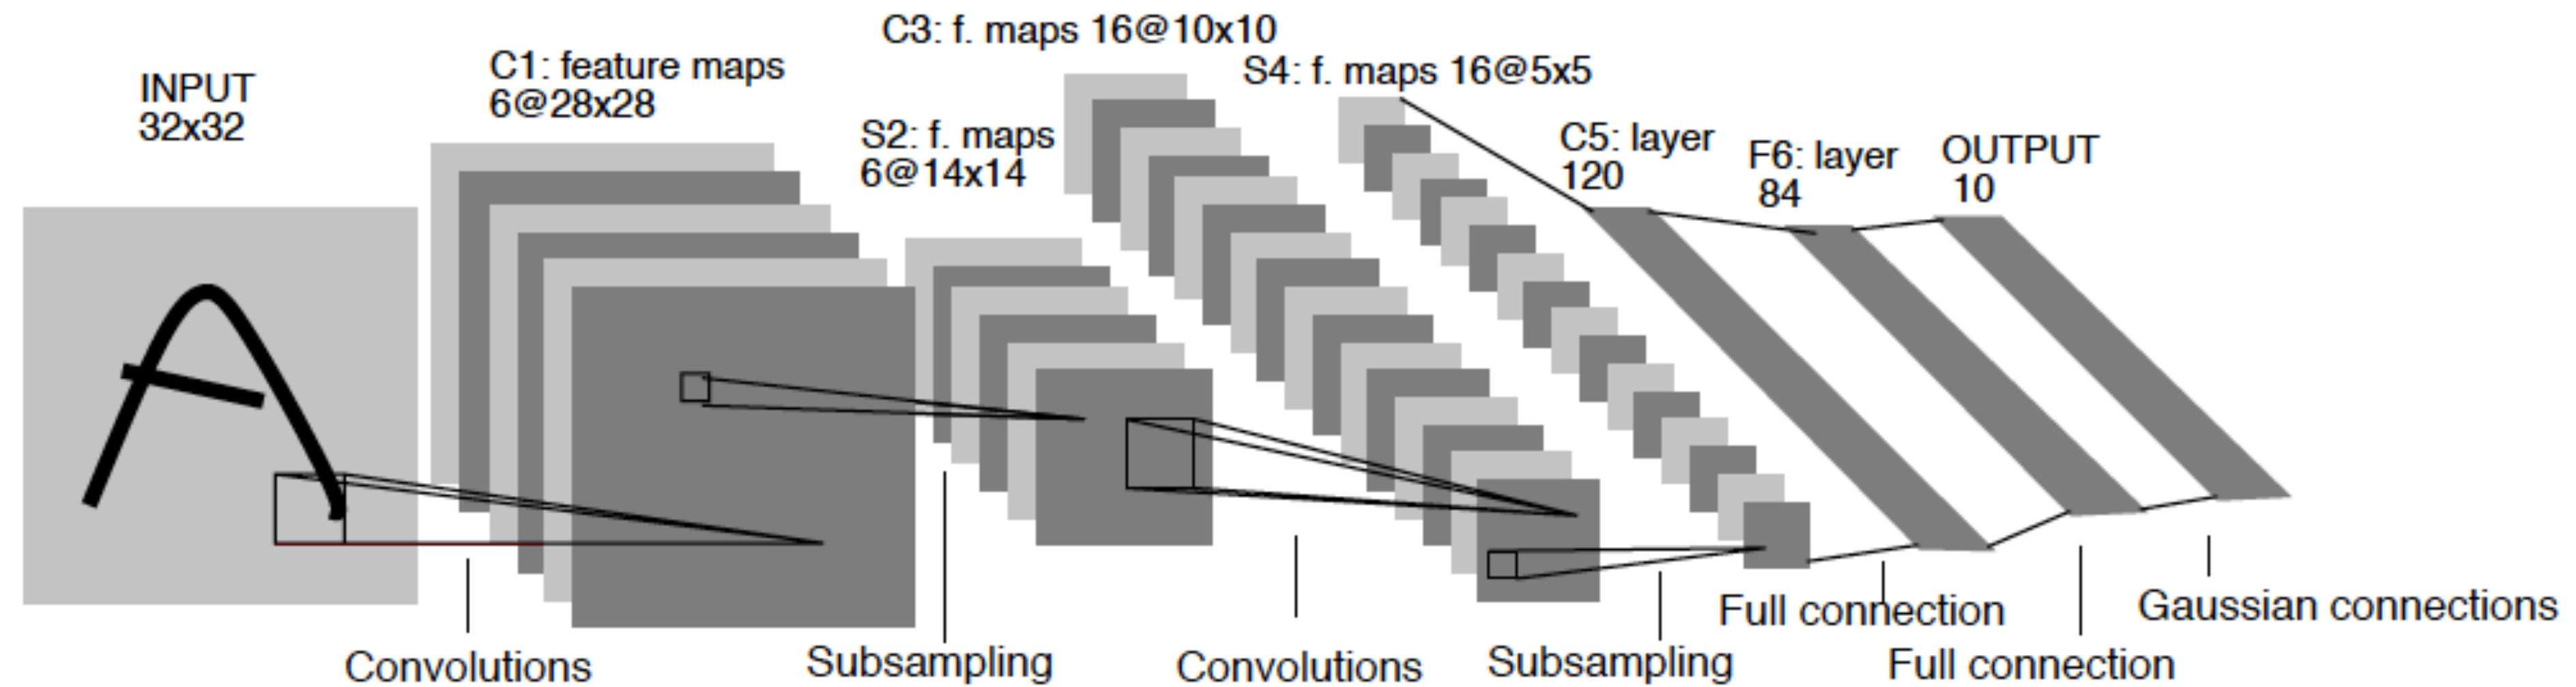
Review of Lecture 12 (Summary)

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires 4 hyperparameters:
 - Number of filters K ,
 - Their spatial extent F ,
 - The stride, S ,
 - And the amount of zero-padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where
 - $W_2 = (W_1 - F + 2P) / S + 1$
 - $H_2 = (H_1 - F + 2P) / S + 1$
 - $D_2 = K$
- With parameter sharing, introduces $F \times F \times D_1$ weights per filters, for a total of $(F \times F \times D_1) \times K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a convolution of the d -th filter over the input volume, with a stride of S , and offset by the d -th bias.

Today's Lecture

- **Convolutional Neural Network Architectures**

LeNet-5



- 5x5 convolutional filters applied at stride of 1
- 2x2 subsampling (pooling) applied at stride of 2
- Architecture is CONV-POOL-CONV-POOL-CONV-FC-FC

Pooling

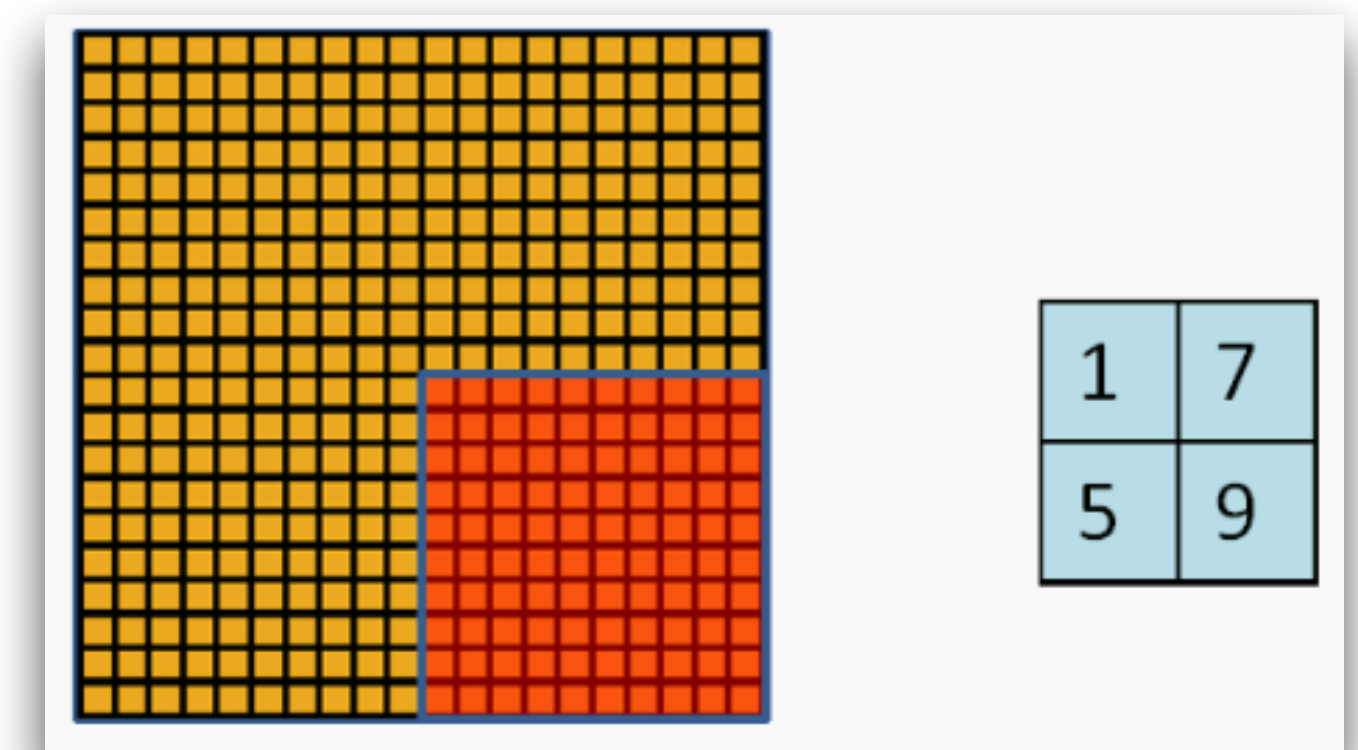

Objective: Achieve spatial invariance by reducing the resolution of activation maps.

- One pooled activation map per map of the previous layer
- Pooling window of $n \times n$ can be of arbitrary size
- Compare subsampling with max pooling:

$$a_j = \text{ReLU}\left(\beta \sum_{N \times N} a_i^{n \times n} + b\right) \quad (\text{subsampling})$$

$$a_j = \max_{N \times N} \left(a_i^{n \times n} u(n, n) \right) \quad (\text{max pooling})$$

windowing function



Pooling

For backpropagation,

- Subsampling layer is treated as usual
- At pooling layer, the error signal is only propagated to the position at:

$$a_j = \arg \max_{N \times N} (a_i^{n \times n} u(n, n)) \quad (\text{results in sparse error maps})$$

Subsampling is clearly more expensive. But is it superior?

No! It's inferior for selecting invariant features and generalizing.

Use max pooling without smoothing or overlap!

AlexNet

Layer architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4


CONV5

MAX POOL3

FC6

FC7

FC8




Details:

- 1000-way softmax classifier
- Convolutional layers, max pooling, ReLU
- SGD with weight decay (batch size=128)
- Dropout on fully-connected layers
- Data augmentation

AlexNet

Input: 227x227x3 images



First layer (CONV1): 96 11x11 filters applied at stride of 4


What is the output volume size? [Hint: $(227-11)/4+1=55$]

Output volume: **55x55x96** Number of parameters?

Parameters: $(11 * 11 * 3)96 + 96 = \mathbf{35K}$

AlexNet

Input: 227x227x3 images
After CONV1: 55x55x96



Second layer (POOL1): 3x3 filters applied at stride 2

What is the output volume size? [Hint: $(55-3)/2+1=27$]

Output volume: **27x27x96** Number of parameters?

Parameters: **0!**


AlexNet

Input: $227 \times 227 \times 3$ images

After CONV1: $55 \times 55 \times 96$

After POOL1: $27 \times 27 \times 96$

...



AlexNet

AlexNet Architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1


[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons

[1000] FC8: 1000 neurons (class outputs)



(Krizhevsky et al. 2012)


Details:

- First use of **ReLU**
- Used Norm layers (across channels - not used anymore)
- Heavy data augmentation
- Dropout 0.5
- Batch size 128
- SGD with **Momentum** 0.9
- LR 1e-2, reduced by 10X manually (val accuracy plateaus)
- L₂ weight decay 5e-4
- 7 CNN **ensemble**: 18.2% error down to 15.4%

ReLU (Rectified Linear Unit)

Computes

$$f(x) = \max(0, x)$$



Pros:

1. Greatly accelerates convergence of SGD (6X) due to linear non-saturating form
2. Inexpensive to compute (threshold the matrix at zero)
3. More biologically plausible

Cons:

1. Can “die” (update weights such that will never activate again)
2. Non-zero centered

Momentum


Definition modifies weight update equation:

$$V_t = \beta V_{t-1} + (1 - \beta) \nabla_w L(W, X, y)$$

$$W = W - \gamma V_t$$

More commonly:

$$V_t = \beta V_{t-1} + \alpha \nabla_w L(W, X, y)$$



Why it works:

- Suppresses noise in the 'right' way
- Dampens oscillations in 'ravines'

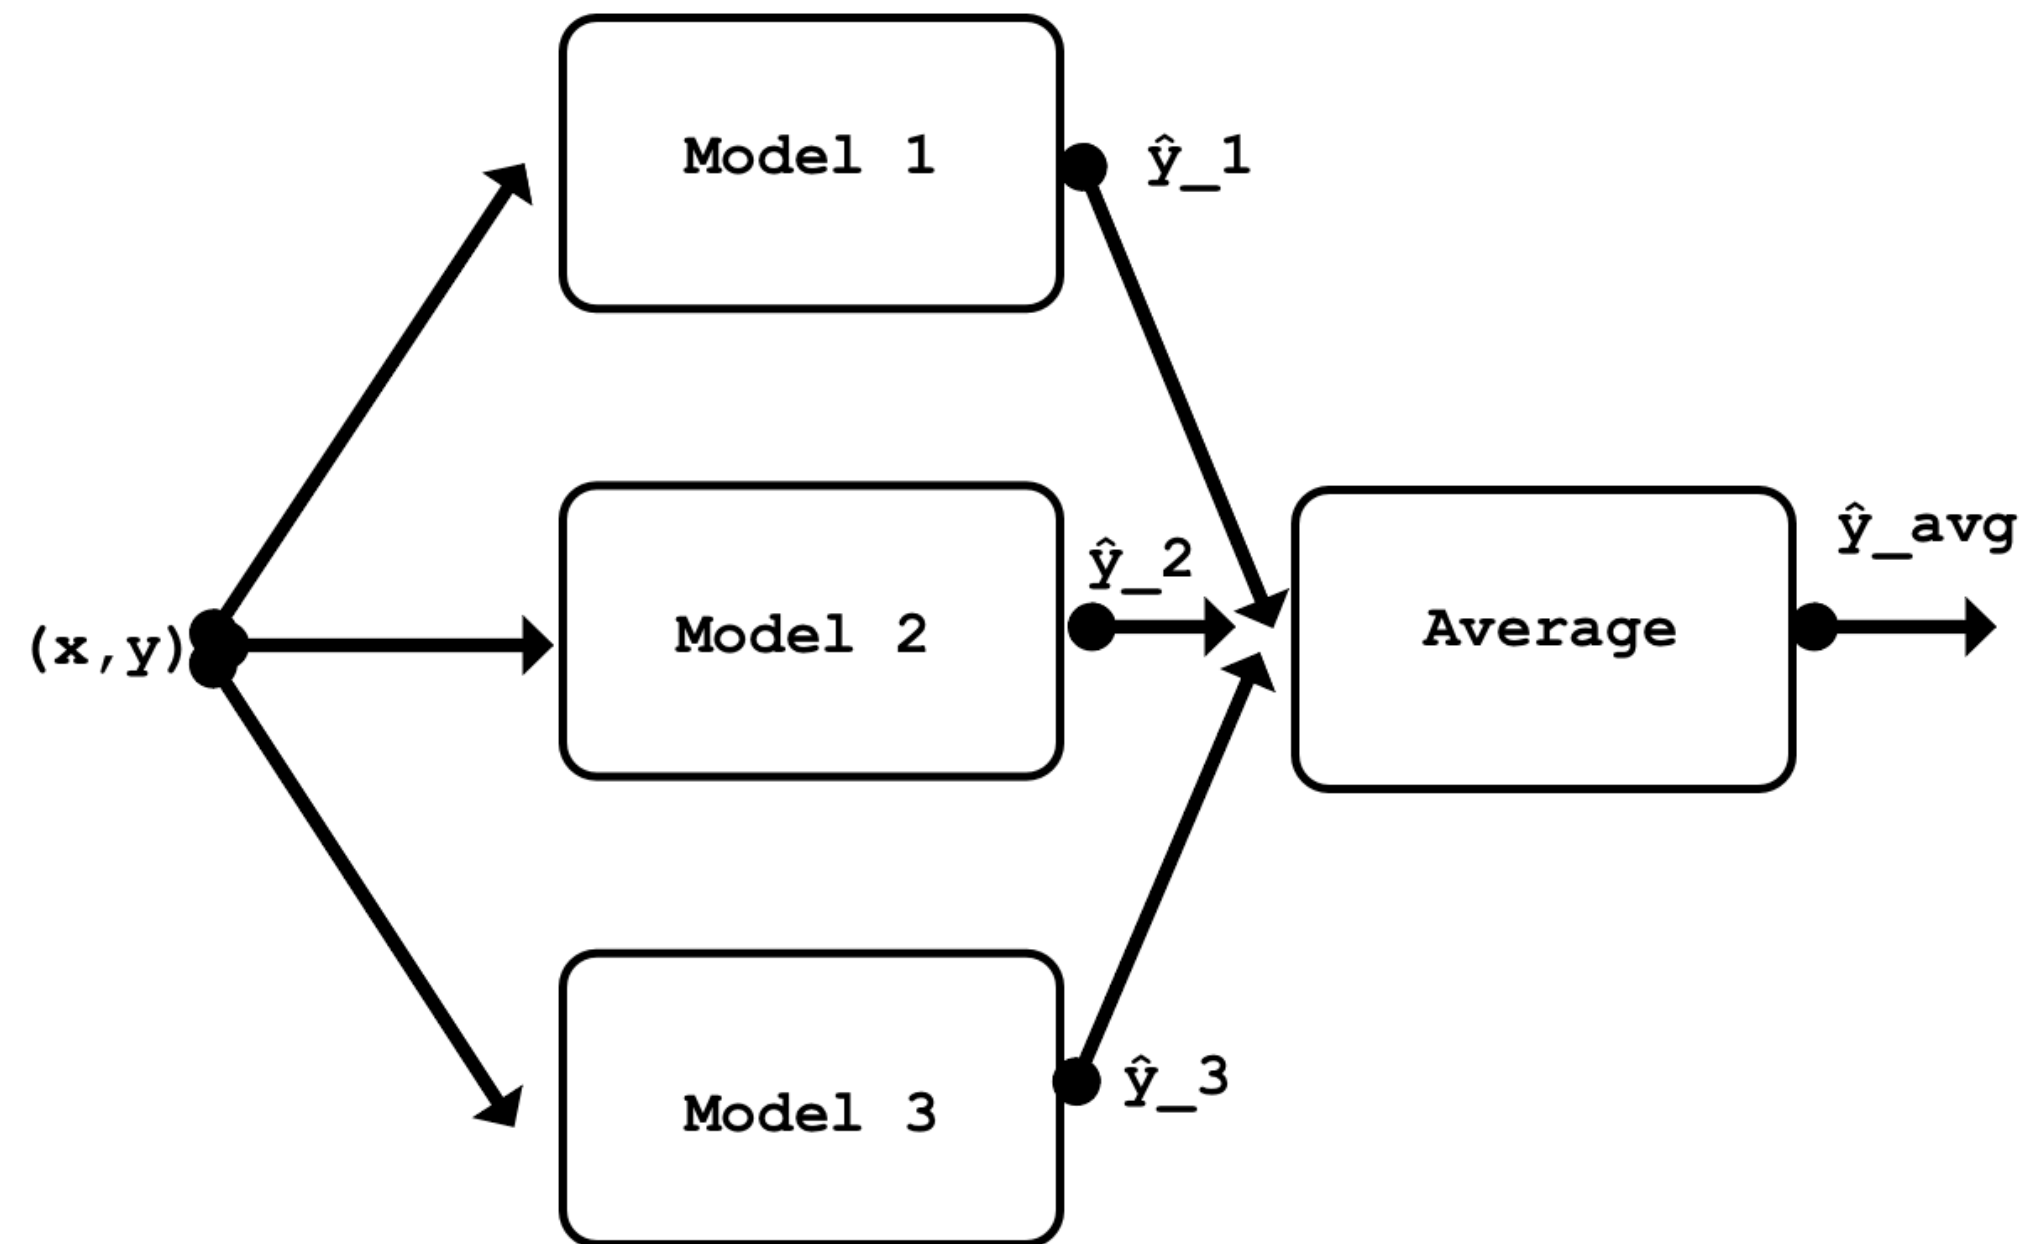
How to use:

- Available with many optimizers (SGD, RMSprop)
- Typical after cross val: $\beta = [0.5, 0.9, 0.95, 0.99]$
- Commonly annealed (start low, end high)

With momentum, the parameter vector will build up “velocity” in any direction with consistent gradient.

Ensembles of Models

Ensemble of Models: Train multiple versions of a model, or multiple independent models.



Approaches may include:

- Same model, different initializations
- Top models discovered during cross val
- Different epoch checkpoints of same model
- Running average of parameters over last few iterations (smoothed version of weights)

Disadvantages? Complexity, computational cost

AlexNet

AlexNet Architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1


[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons

[1000] FC8: 1000 neurons (class outputs)



(Krizhevsky et al. 2012)

[55x55x48] x 2!

- Trained on GTX 580 GPU (3 GB memory)
- Network split across 2 GPUs
- Half the feature maps per GPU

AlexNet

AlexNet Architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1


[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons

[1000] FC8: 1000 neurons (class outputs)



(Krizhevsky et al. 2012)

CONV1, CONV2, CONV4, CONV5 connect only with activation maps on same GPU.

AlexNet

AlexNet Architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1


[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons


[1000] FC8: 1000 neurons (class outputs)




(Krizhevsky et al. 2012)

CONV3, FC6, FC7, and FC8 connect with all activation maps in preceding layer (communicate across GPUs).

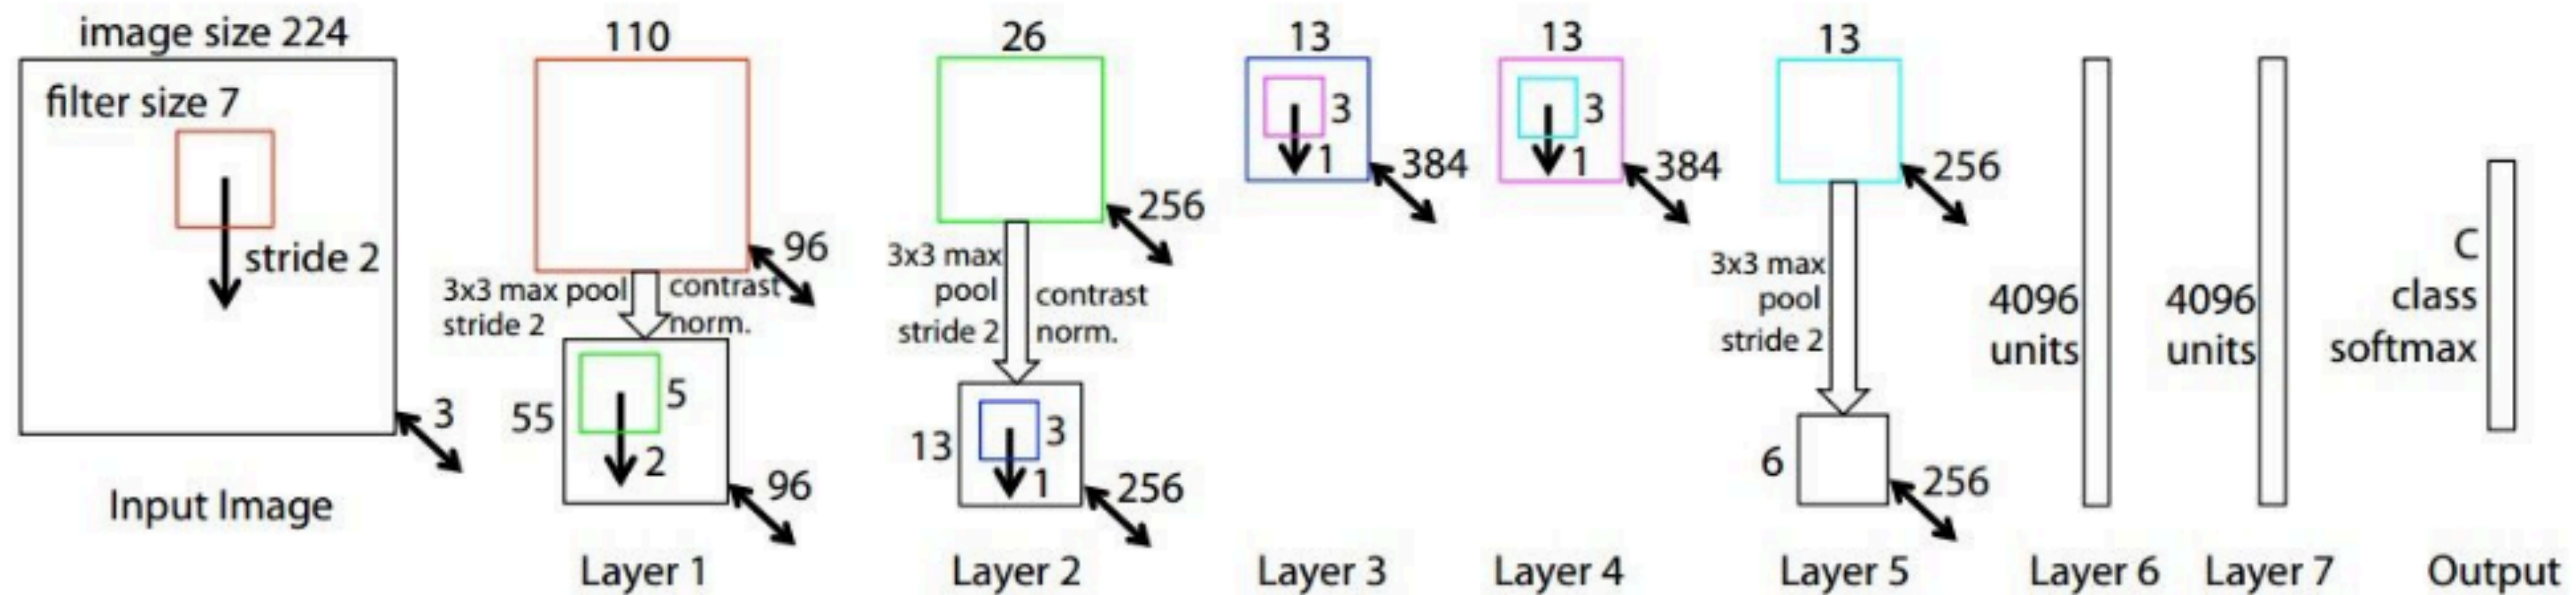
AlexNet and ImageNet (ILSVRC contest)



AlexNet and ImageNet (ILSVRC contest)



ZFNet



You get the idea!


Like AlexNet, but:

- CONV1 changed from 11x11 with stride 4 to 7x7 with stride 2
- CONV3,4,5 changed from 384, 384, and 256 filters to 512, 1024, and 512 filters

ImageNet top-5 error improved from 16.4% to 11.7%

(Zeiler and Fergus 2013)

AlexNet and ImageNet (ILSVRC contest)



VGGNet

Smaller filters, deeper networks

8 layers (AlexNet)




16-19 layers (VGG16Net)

Smallest that looks at neighbors!

Only 3x3 CONV stride 1, pad 1
and 2x2 POOL stride 2

**ImageNet top-5 error improved
from 11.7% to 7.3%**



AlexNet



VGG16

VGG19

(Simonyan and Zisserman 2014)

VGGNet

Smaller filters, deeper networks

Why use smaller filters (e.g., 3x3 conv)?

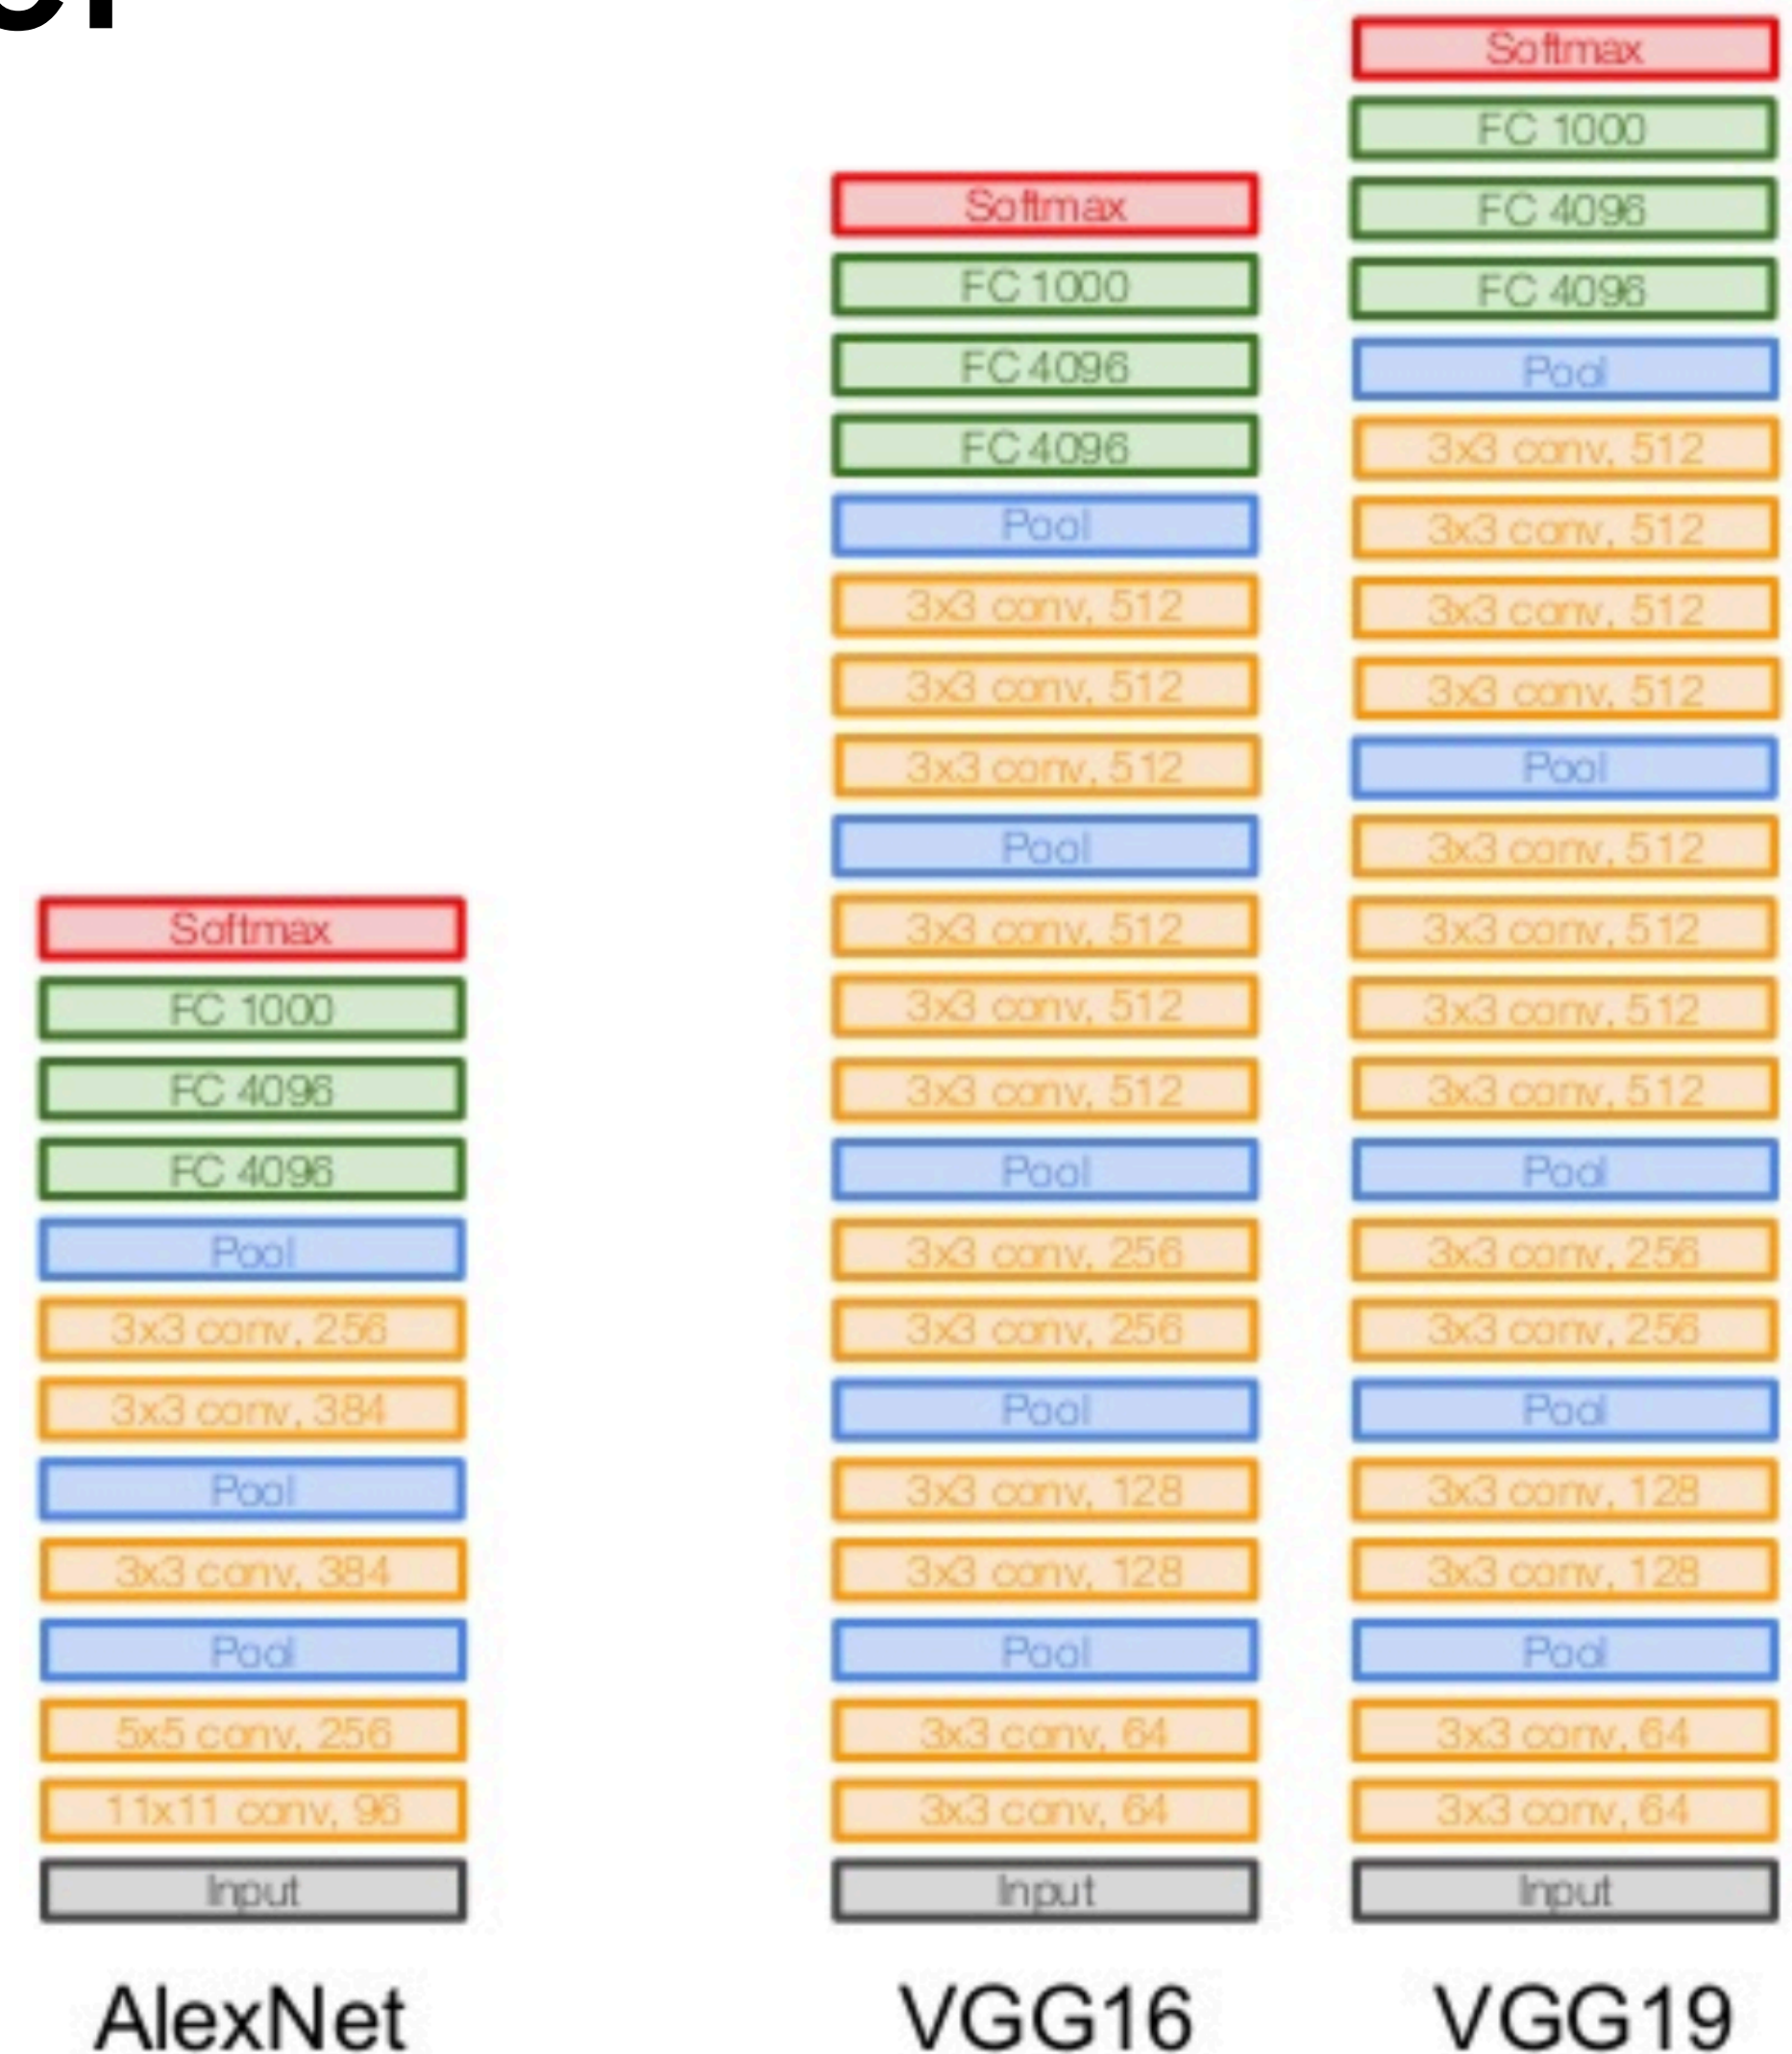
Related: What is the receptive field of three 3x3 conv layers (stride 1)?

Same **effective receptive field** as one 7x7 conv layer...

... but deeper (more non-linearities) ...

... and fewer parameters!

$3 * (3^2 C^2)$ versus $7^2 C^2$ for C channels per layer.



(Simonyan and Zisserman 2014)

VGG16 memory usage and parameters

INPUT: [224x224x3] memory: $224*224*3=150\text{K}$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2\text{M}$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2\text{M}$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800\text{K}$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6\text{M}$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6\text{M}$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400\text{K}$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200\text{K}$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100\text{K}$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25\text{K}$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

Most memory is used
in early conv layers

Most parameters used
in FC layers

Total memory: $24\text{M} * 4$ bytes or approx 96 MB per image (forward pass; 2X to include backward!)

Total parameters: 138M!!

VGG16 memory usage and parameters

INPUT: [224x224x3] memory: $224*224*3=150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800K$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400K$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200K$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100K$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$



VGG16

**Common
names**

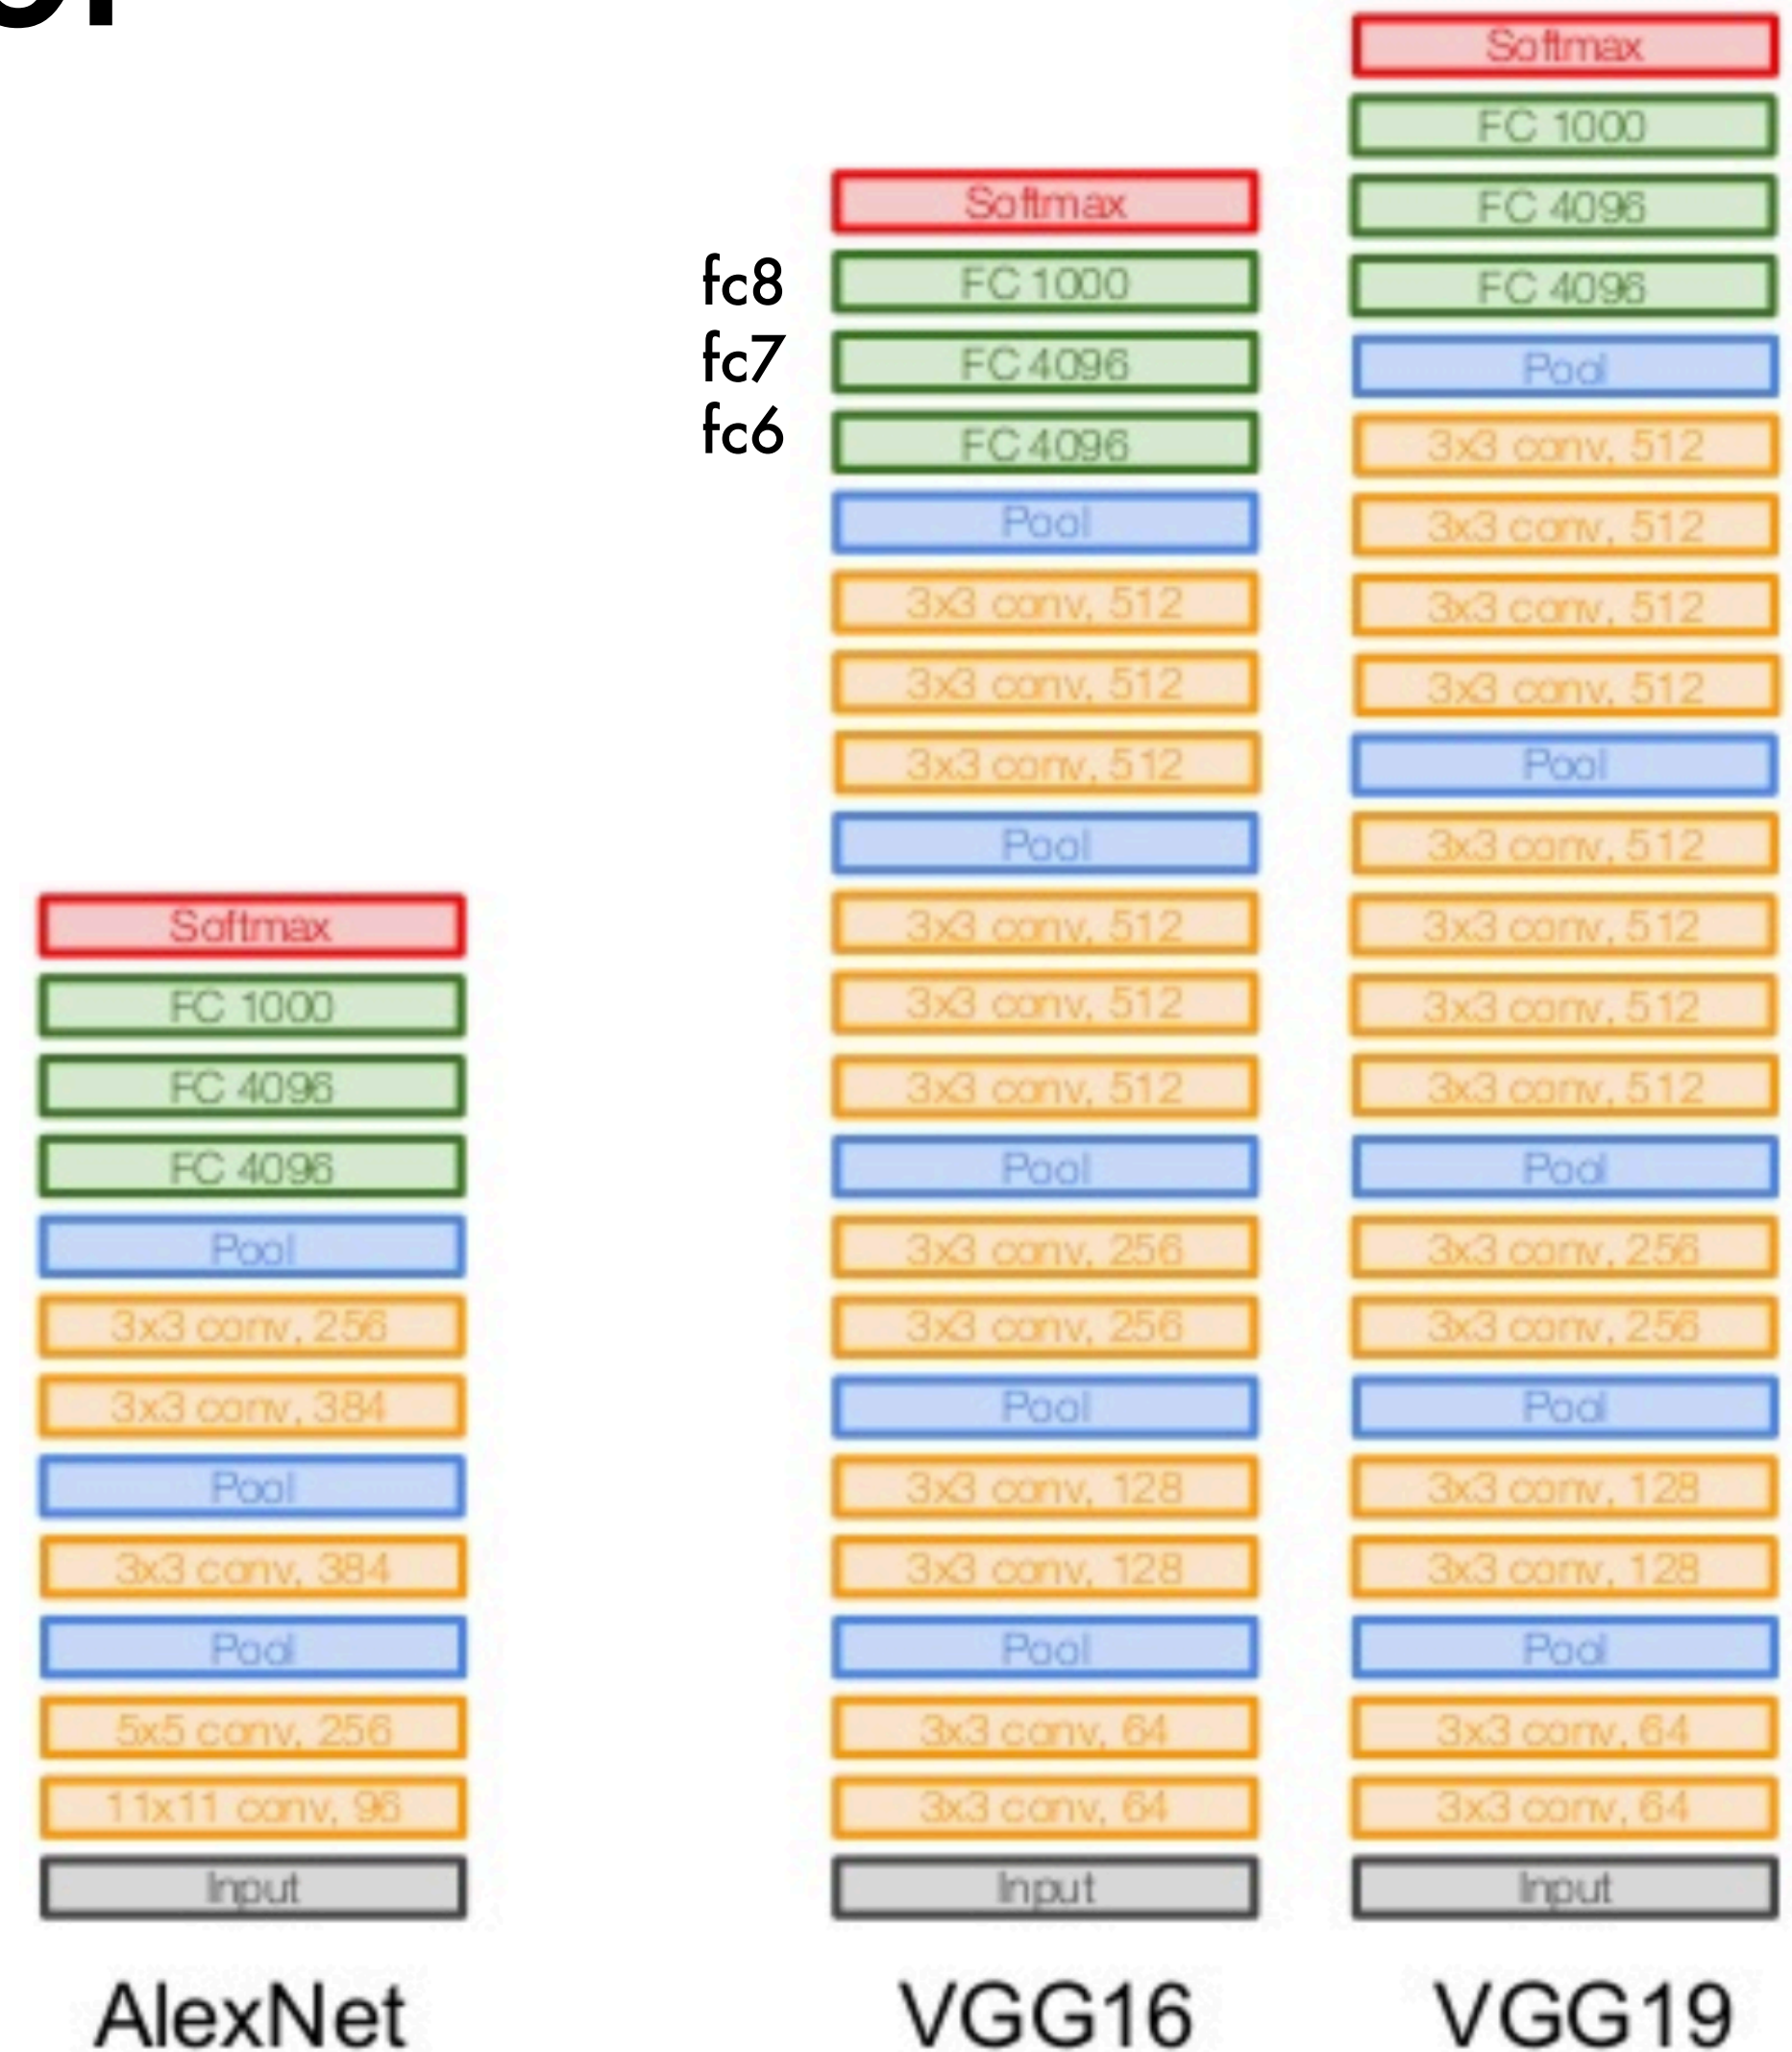
Total memory: $24M * 4$ bytes or approx 96 MB per image (forward pass; 2X to include backward!)

Total parameters: 138M!!

VGGNet

Additional notes:

- 2nd in classification, 1st in localization in ILSVRC'14
- Similar training as AlexNet
- Didn't use the NORM
- VGG19 slightly better (more memory)
- Use ensembles for best results
- fc7 features generalize well to other tasks

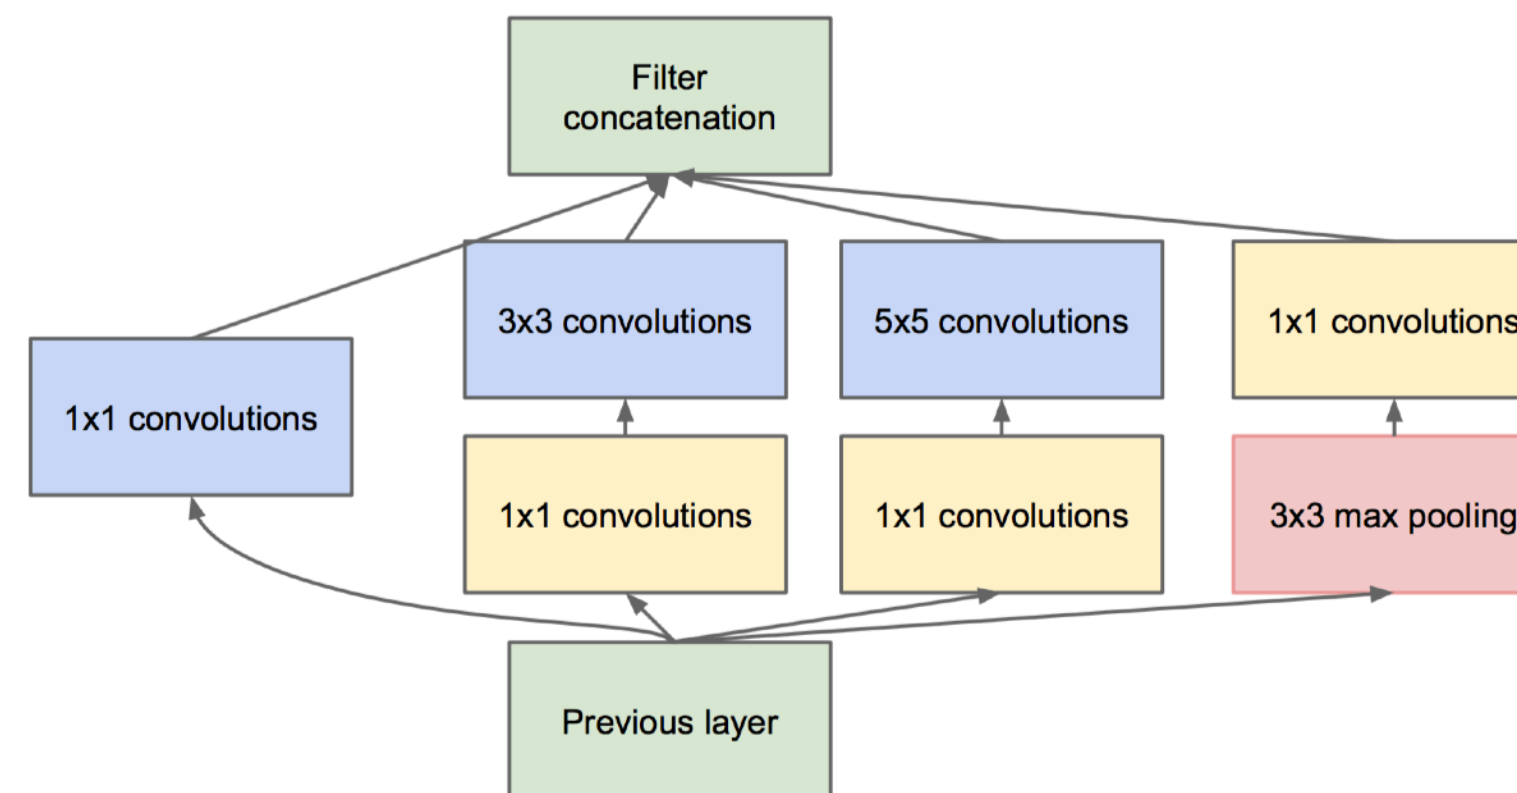


(Simonyan and Zisserman 2014)

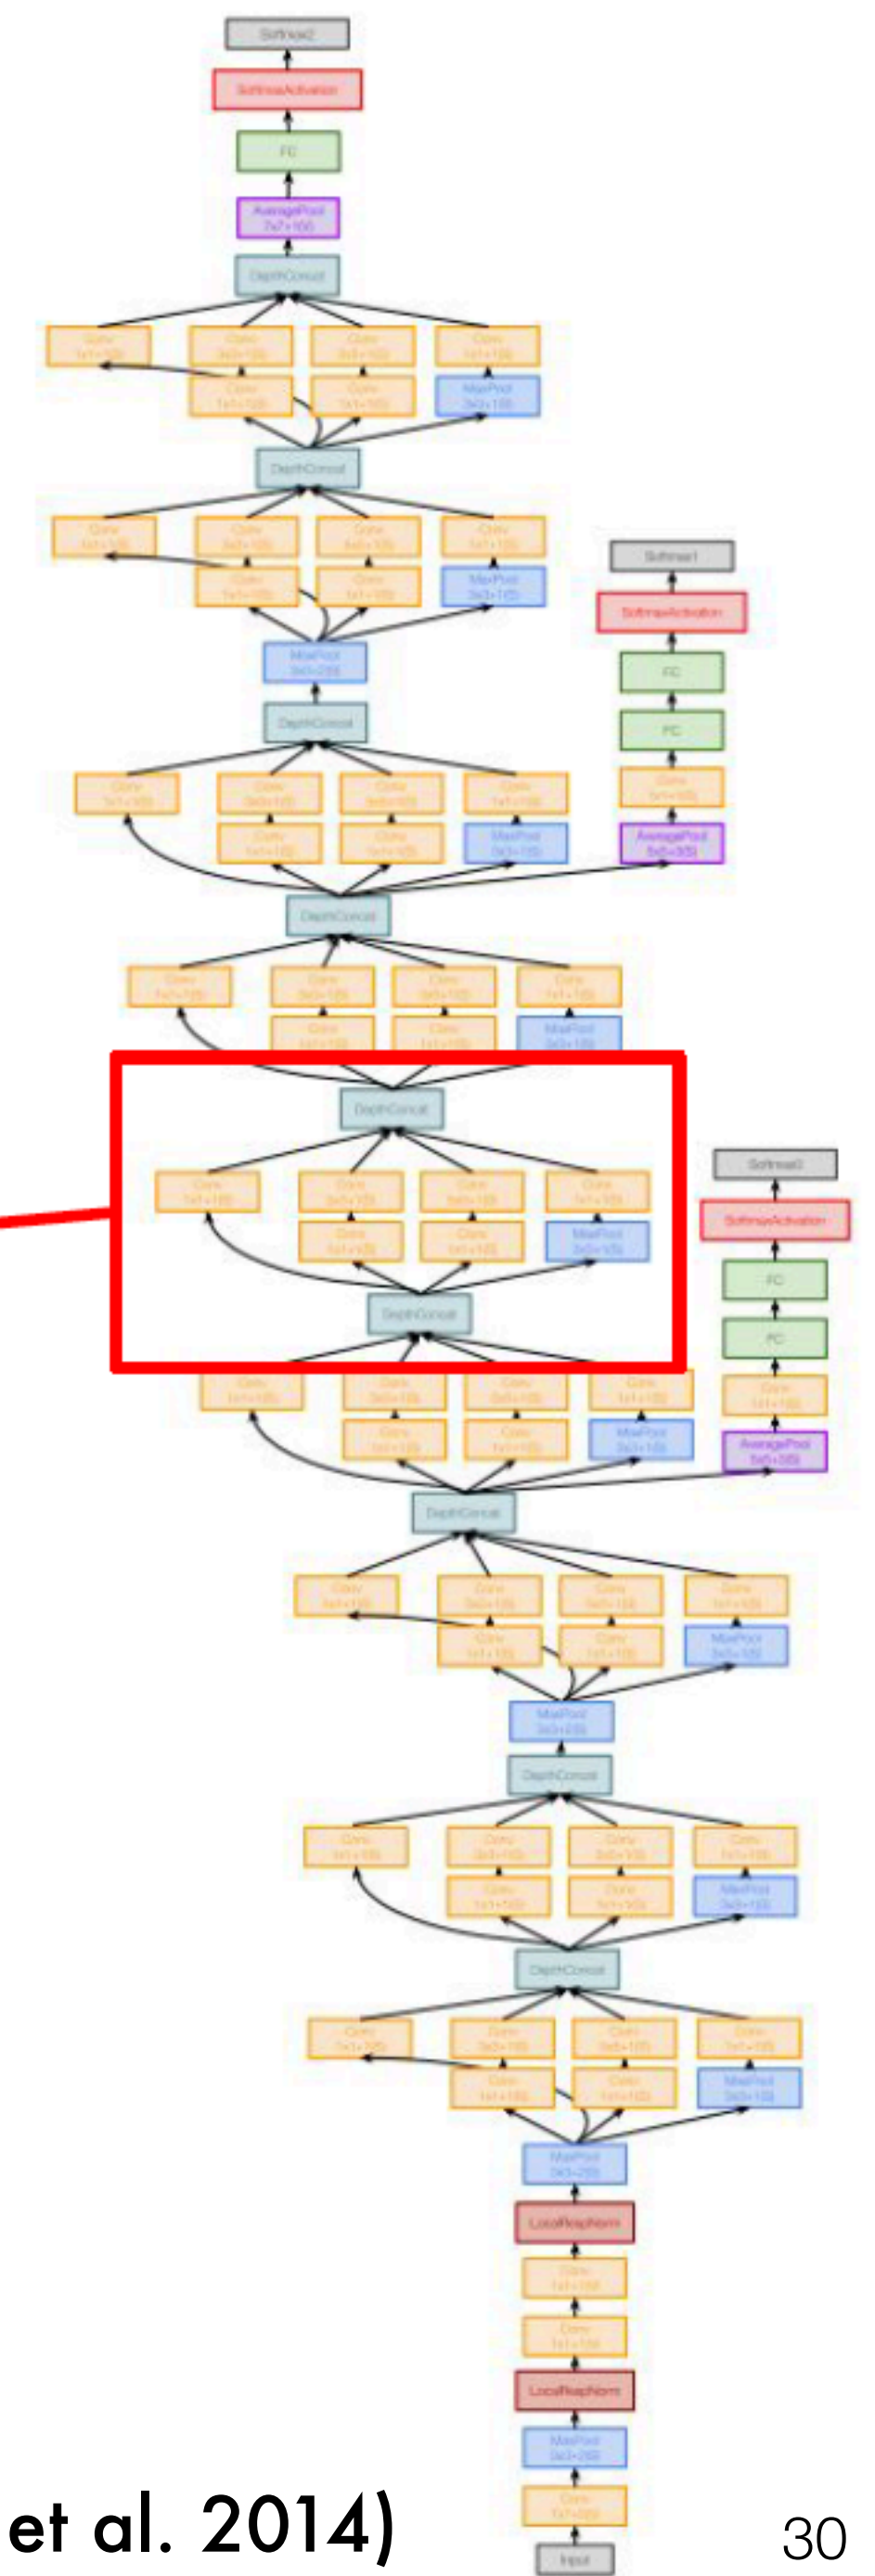
GoogLeNet

**Deeper networks,
with computational efficiency**

- 22 layers
- Efficient 'inception' module
- No FC layers!
- Only 5 million parameters!
(1/12th of AlexNet)
- ILSVRC'14 classification winner
(6.7% top-5 error)



Inception module




(Szegedy et al. 2014)

GoogLeNet

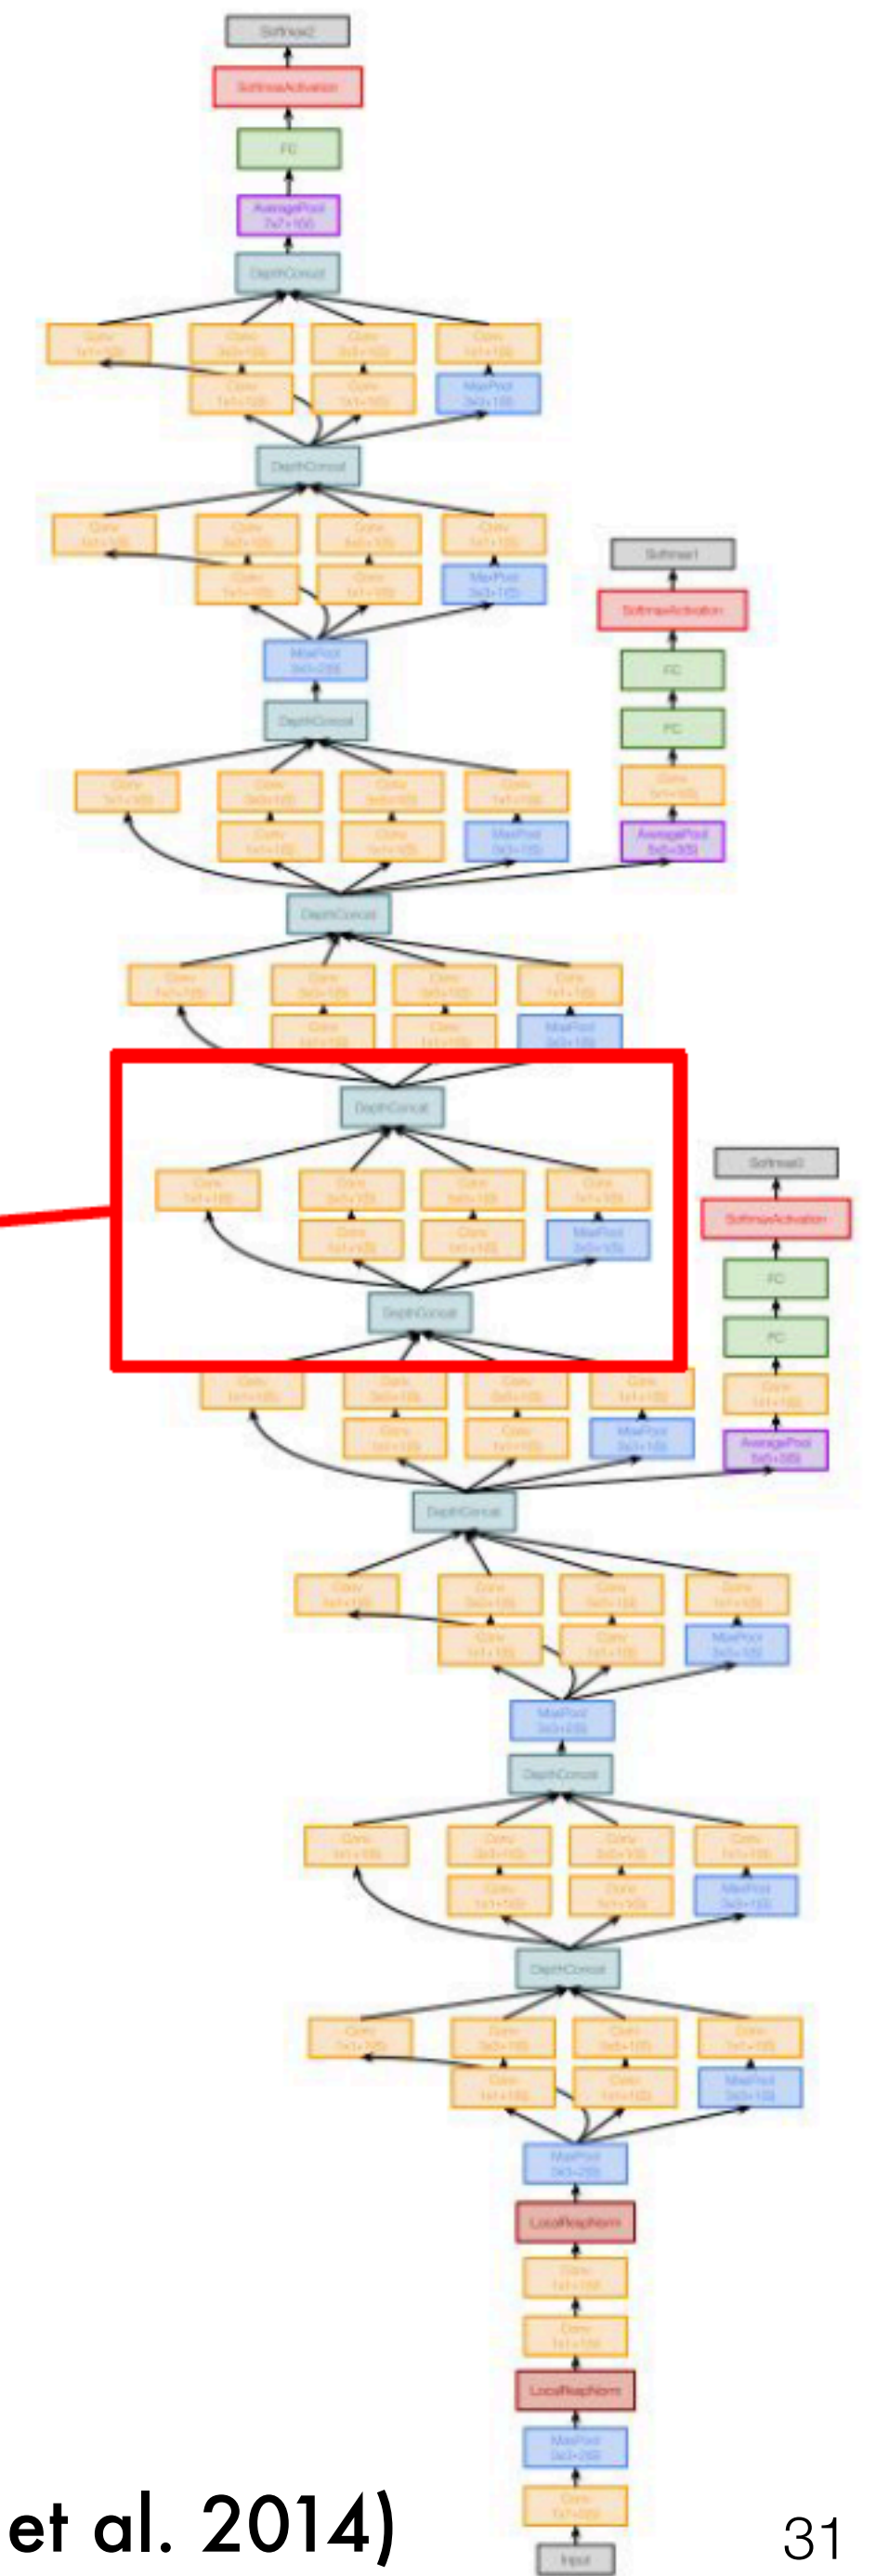
**Deeper networks,
with computational efficiency**

‘Inception model’

Design a good local network topology (network within a network) and then stack these modules on top of each other.

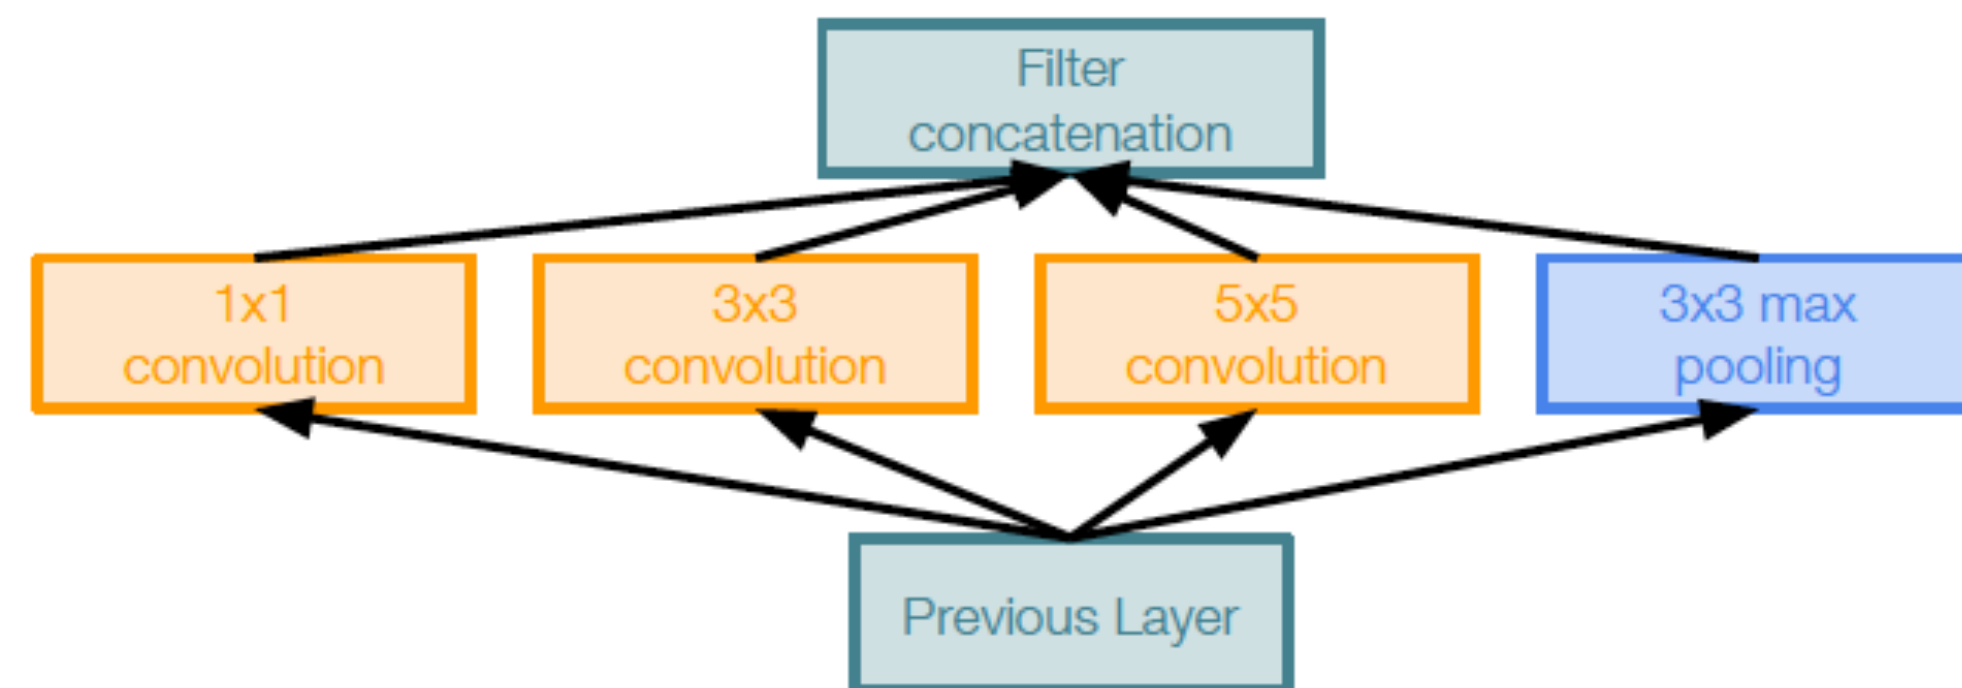


Inception module



(Szegedy et al. 2014)

GoogLeNet



Naive Inception module

Apply parallel filter operations on the input from previous layer


- Multiple receptive field sizes for convolution (1x1, 3x3, 5x5)
- Pooling operation (3x3)

Concatenate all filter outputs together depth-wise

What is the problem with this?
[Hint: Computational complexity!]

GoogLeNet

$$28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672!$$



Naive Inception module

Module input: $28 \times 28 \times 256$

What is the problem with this?
[Hint: Computational complexity!]

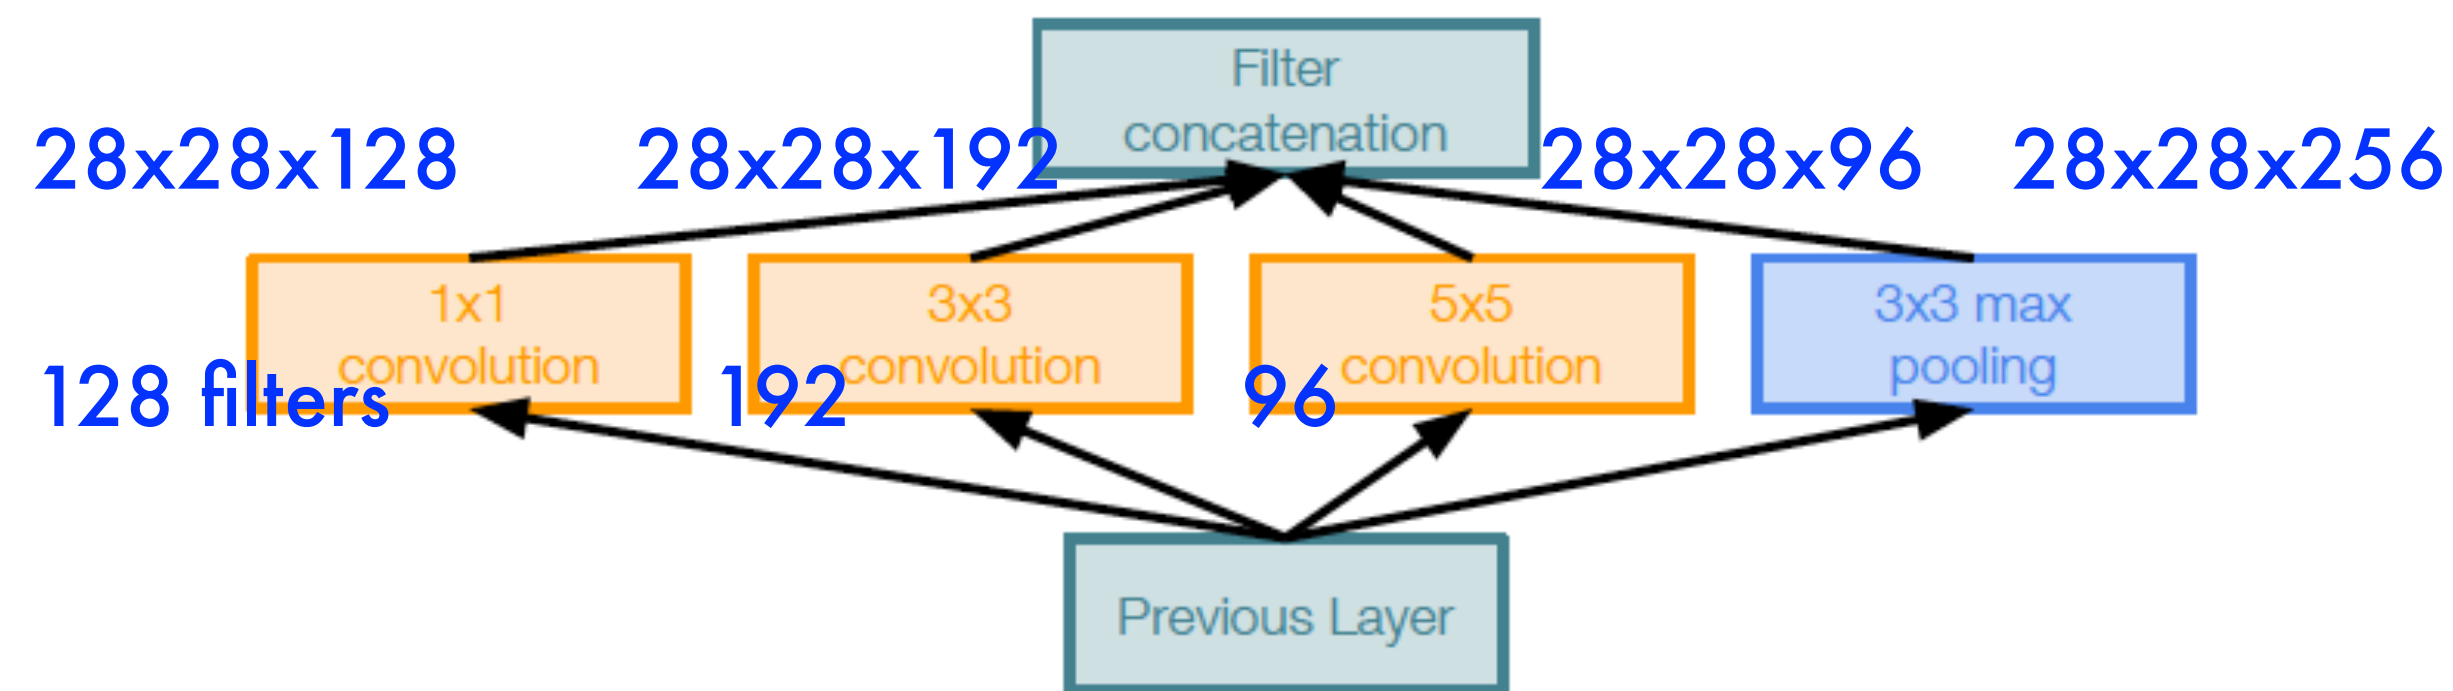
What is the output size of the 1×1 conv with 128 filters?

What is the output size of all of the different filter operations?

What is the output size after filter concatenation?

GoogLeNet

$$28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672!$$



Naive Inception module

Conv Ops:

[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$

[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

Total of 854M ops!

Very expensive compute!

Pooling layer also preserves feature depth, which means total depth after concatenation can only grow at every layer!

See next lecture for solution!

Further reading

- Abu-Mostafa, Y. S., Magdon-Ismail, M., Lin, H.-T. (2012) Learning from data. [AMLbook.com](http://www.amlbook.com).
- Goodfellow et al. (2016) Deep Learning. <https://www.deeplearningbook.org/>
- Boyd, S., and Vandenberghe, L. (2018) Introduction to Applied Linear Algebra - Vectors, Matrices, and Least Squares. <http://vmls-book.stanford.edu/>
- VanderPlas, J. (2016) Python Data Science Handbook. <https://jakevdp.github.io/PythonDataScienceHandbook/>